

Probability and Mathematical Statistics: Final Project

Due on December 31, 2022 at 11:59am

Name: **Ren Hui, Wanchen Su**
Student ID: 2021533089, 2021533067

Problem 1 (Part 1: classical Bandit Algorithms)**Solution**

1,2. See jupyter notebook.

3. According to python simulation, the results are:

results for ϵ -greedy:

	$\epsilon = 0.1$	$\epsilon = 0.5$	$\epsilon = 0.9$
reward	3415.895	3083.32	2748.985
θ	[0.699565 0.501675 0.398460]	[0.700644 0.496887 0.398588]	[0.700090 0.499850 0.398113]

results for UCB:

	$c = 0.1$	$c = 0.5$	$c = 0.9$
reward	3410.44	2974.035	2824.62
θ	[0.700882 0.492749 0.384489]	[0.700674 0.499007 0.398861]	[0.701356 0.498778 0.400231]

results for TS:

	a,b = [[1, 1], [1, 1], [1, 1]]	a,b = [[601, 401], [401, 601], [2, 3]]
reward	3481.33	3490.945
θ	[0.6995964 0.45887541 0.37554019]	[0.68296885 0.4001996 0.362563]

4. Oracle value:

$$E(r) = 5000 \times 0.7 = 3500$$

gap of ϵ -greedy:

	$\epsilon = 0.1$	$\epsilon = 0.5$	$\epsilon = 0.9$
reward	3415.895	3083.32	2748.985
gap	84.105	416.68	751.015

gap of UCB:

	$c = 0.1$	$c = 0.5$	$c = 0.9$
reward	3410.44	2974.035	2824.62
gap	89.56	525.965	675.38

gap of TS:

	a,b = [[1, 1], [1, 1], [1, 1]]	a,b = [[601, 401], [401, 601], [2, 3]]
reward	3481.33	3490.945
gap	18.67	9.905

As is shown above, with the given parameters, TS Algorithm is the best.

In the ϵ -greedy algorithm, ϵ decides the probability of choosing the max estimate of all the arms or choosing a random arm. Which means the larger ϵ is, the more evenly spread are the tests.

In the UCB algorithm, the parameter c balances the estimation of θ_j , and considers the number of trials on that arm. Increasing c means considering more about the number of test on that arm, which will allow the decision maker(DM) choose some of the less experimented arms. This helps in instances where the number of tests on a arm is quite small and therefore the θ of that arm is way less than it

should be.

If the initial value of α_j, β_j we passed in is too small, then the result of the first few tests may influence the final result greatly, and only a few dozen tests are on the second and third arm. If this happens, the estimate value of θ will be rather inaccurate. At the same time, the aggregated reward will be quite large because the DM spent very little time exploring. On the other hand, if the initial value of α_j, β_j we passed in is relatively large, this value of a and b may influence the final result greatly. As in the trial with parameter $[[601, 401], [401, 601], [2, 3]]$, we can see that the final estimation of arm two is very close to $\frac{401}{401+601}$. In this case, the parameter we passed in is correct in which is the arm with the largest oracle value, so the aggregated reward is quite large. However, if the parameter passed in is contradictory to the actual oracle, it will take quite a long time to adjust. Consequently the aggregated reward is smaller.

More impacts of ϵ , c and α_j, β_j concerning exploration-exploitation will be discussed in 5.

5. In the case of all algorithms, more exploitation means that the total reward we get is larger, but if the estimate of each arm differs too much from the oracle value, after all tests, the total aggregated reward becomes smaller. More exploration means the estimated value of the reward of each arm is more accurate, it helps the decision maker choose the arm better, but the exploration means that the rewards during exploration is less.

For the ϵ -greedy algorithm, the larger ϵ is, there is more chance of exploration than exploitation. So while ϵ grows, the gaps between the algorithm estimate and the oracle value decreases, but the sum of the reward of all trials is smaller because the DM spent too much time exploring when it is not that necessary.

As for the UCB Algorithm, the exploration-exploitation trade-off depends on the value of c . Similar to the ϵ -greedy algorithm, the larger c is, there is more exploration and less exploitation. As can be deduced from the data in the python simulation, for larger c , the estimated value for the third arm is more accurate. This means that there are more times when the decision maker choose arm c . Consequently, the aggregated rewards over N time slots are less.

The TS Algorithm is rather different from the previous ones. Its exploration is very limited. The arm selected is always the arm with the greatest current reward estimation. So arms with small oracle values may end up being pulled only a few times. Which makes the estimate value of arms 2 and 3 rather inaccurate. But the aggregated reward gained over the N trials should be the largest of all the algorithms. TS performs best in this case because the parameters a and b we passed in corresponds with the real oracle value. On the other hand, as the example in 5, if the parameters passed in is inaccurate, it could have a large influence on the aggregated reward. This is a default for having too little exploration.

6. We searched the internet for more information on depend arm bandit problem, and used the thought of the c-UCB algorithm in a paper(see part 1 reference reading 3) to improve the classical bandit algorithms.
 - a. Data choice: After research, we found that the dependent multiarm bandit problem is often used in machine learning to decide which items to promote for each user. So besides using computer generated data to test our dependent bandit algorithm, we also used authentic data downloaded from the internet. We considered the real-life application of the dependent multi-arm bandit problem, more specifically, the problem of movie promotion. For app designers, they want to push book advertisements according to user preference, and user preference is linked to certain traits of the user in question. We downloaded a dataset of viewer ratings on movies. The choice of movie we push for the viewer is the arm we choose, and the viewer's final rating is the reward. To match a three arm bandit machine, we first implemented the dependent algorithm on choosing three movies to promote. Afterwards we also tried the broader version of 18 movies.

- b. Performance evaluation: To make the result more clear, we use expected regret ($E[Reg(T)]$) to analyze the performance of our algorithm. The regret of an algorithm is defined as the difference between the cumulative reward of a genie policy, that always pulls the optimal arm, and the cumulative reward obtained by the algorithm over T rounds. So obviously its expectation in each round is the difference between the mean reward of the best arm and the mean reward of the chosen arm.
- c. Dependency: After processing the data, we obtained pseudo-rewards, which provides an upper bound on the conditional expectation of an arm l given that we have received a reward r at another arm k . We use the definition of pseudo-rewards from the paper "Multi-Armed Bandits with Correlated Arms", which is: "Suppose we pull arm k and observe reward r , then the pseudo-reward of arm l with respect to arm k , denoted by $s_{l,k}(r)$, is an upper bound on the conditional expected reward of arm l , i.e.,

$$E[R_l | R_k = r] \leq s_{l,k}(r)$$

without loss of generality, we define $s_{l,l}(r) = r$." This information can be obtained in practice through either domain/expert knowledge or from controlled surveys. Similarly, define $\phi_{l,k} = E[s_{l,k}(R_k)]$.

- d. Dependent bandit algorithm: We realized both the dependent UCB algorithm and the dependent TS algorithm. Since both are similar in principle, we will mainly discuss the dependent UCB algorithm.

The basic idea is to use dependency to eliminate some of the less competitive arms, so there are less choices each round, and the choices are closer to the best choice.

Step 1: Identify the set of significant arms The set of significant arms at round t S_t is defined to be the set of arms that have at least $t/arms$ samples. In other words, S_t is the set of arms that have relatively large number of samples. So we will use the data of these arms to eliminate less competitive arms by dependency.

Step 2: Identify the set of competitive arms First, find the largest estimated theta value so far by UCB(or TS) algorithm. Let that arm or the few arms that have the best theta be C_t . Then, for each arm l , find the smallest $\phi_{l,k}$ with $k \in S_t$. If this value is less than the largest estimated theta value so far, then this arm l is not competitive. Note that the ϕ value changes at each round so arm l is not eliminated forever. Denote the set of all competitive arms to be A_t .

Step 3: Use bandit algorithm on $A_t \cup C_t$ According to the steps above, $A_t \cup C_t$ is the set of all competitive arms at round t .

The specific algorithm we used is shown below.

```

Input: pseudo-rewards  $s_{l,k}(r)$ 
Initialize:  $count[k] = 0, theta[k] = \infty$  for all  $k$  in arms
for each round  $t$  do
    Find  $S_t$ , the arm(s) that have been pulled significant times till the  $t-1$  round.
    Define  $k^{emp}(t) = \operatorname{argmax}_{k \in S_t} \theta_k(t)$ 
    Initialize the competitive set  $A_t$  as an empty set.
    for  $k \in arms$  do
        if  $\min_{l \in S_t} \phi_{k,l}(t) \geq \theta_{k^{emp}(t)}$  then
            add arm  $k$  to  $A_t$ 
        end
    end
    Apply UCB to  $A_t \cup C_t$ 
    Receive reward and update reward,  $count[k]$ ,  $theta[k]$ , pseudo-reward.
end

```

Algorithm 1: Dependent UCB

*7. reference reading:

- 1 <https://zhuanlan.zhihu.com/p/84140092>, <https://zhuanlan.zhihu.com/p/84200578>,
<https://zhuanlan.zhihu.com/p/84338172>
- 2 Rahul Singh, Fang Liu, Yin Sun, Ness Shroff, *Multi-Armed Bandits with Dependent Arms*
- 3 Samarth Gupta, Shreyas Chaudhari, Gauri Joshi, Osman Yağın, *Multi-Armed Bandits with Correlated Arms*
- 4 S. Agrawal and N. Goyal, “Further optimal regret bounds for thompson sampling,” in *Artificial Intelligence and Statistics*, pp. 99-107, 2013.

Problem 2 (Part 2: Bayesian Bandit Algorithms)**Solution**

1. See jupyter notebook for the simulation process. We simulated with the oracle value : **simulations:**

parameter	θ_1	θ_2	λ
test 1	0.7	0.3	0.9
test 2	0.4	0.5	0.9
test 3	0.1	0.3	0.9

Let $r(n)$ be the reward over n pulls. **simulations:**

result	always choose best $E(r(25))$	simulation $r(25)$	always choose worst $E(r(n))$
test 1	6.497471408615705	5.758137590942908	2.784630603692444
test 2	4.641051006154074	4.247677424676706	3.71284080492326
test 3	2.784630603692444	2.2136432344092114	0.928210201230815

2. If both θ_1 and θ_2 are relatively large, then we could end up with always pulling the first arm that succeeded. For example, if $\theta_1 = 0.8$, $\theta_2 = 0.9$, and the first arm we pulled is the first one, then there is a large possibility that the DM will always pull arm 1. We used 0.8 and 0.9 as the oracle value and asked the python simulation to print out the arm it chooses each time, and it shows clearly that it almost always sticks with the arm it chose in the first round. In fact, as it turns out, if the two arms have close possibility of giving a reward, this policy tend to behave not that well.

3. Story proof:

Let us be the decision maker(DM) of a two-armed bandit problem. Presume that the experiment lasts forever, but the reward returned each time is reduced by multiplying λ , where $0 < \lambda < 1$. Put in mind that we as DM, do not know the actual success probability of each arm. So let $R(\alpha_1, \beta_1, \alpha_2, \beta_2)$ represent the expected reward given the previous success and failures. So according to LOTE, the expected reward of pulling the first arm is:

$$R_1(\alpha_1, \beta_1) = P(\text{success}|\alpha_1, \beta_1) \times (\text{reward if succeed}) + P(\text{failure}|\alpha_1, \beta_1) \times (\text{reward if failure})$$

Which is:

$$R_1(\alpha_1, \beta_1) = \frac{\alpha_1}{\alpha_1 + \alpha_2} (1 + \lambda R(\alpha_1 + 1, \beta_1, \alpha_2, \beta_2)) + \frac{\alpha_2}{\alpha_1 + \alpha_2} (\lambda R(\alpha_1, \beta_1 + 1, \alpha_2, \beta_2))$$

R_2 is similar, the only difference is in the case of arm 2 we should change α_2 and β_2 accordingly. Since we will always make the best decision at the moment, $R(\alpha_1, \beta_1, \alpha_2, \beta_2) = \max(R_1(\alpha_1, \beta_1), R_2(\alpha_2, \beta_2))$. Therefore, the given recurrence equation holds.

4. The above recurrence equation can be solved by using dynamic programming in python simulation. We only need to type in the equation and manually set a ending condition, such as count the number of times it has been called repeatedly, and if it exceeds a certain n , then the function will return $\frac{\alpha}{\alpha+\beta}$. See jupyter notebook part2.ipynb for the simulation.
5. The optimal policy will be to alter the algorithm in 4 slightly, so that after solving the equation with dynamic programming, we will get a 4d matrix, storing the corresponding decision we should make with the known previous results. It should be noted that in this step, the case where either arm can be chosen must be threatened with care. We denoted it as 0, rather than arm1 or 2, and when the DM sees a 0 in the matrix it randomly chooses an arm to pull every trial. In this way the matrix will not effect the DM more than it should. The DM will only need to check the matrix to know the arm

it should choose. The simulation for the improved strategy is in part2.ipynb. The result is as below:
simulations:(results are averaged over 200 trials)

test	actual theta	intuitive policy	Bayesian policy
test 1	0.8 0.3	6.526969462197445	6.768196035314434
test 2	0.5 0.7	5.889886701348879	5.964659433792904
test 3	0.5 0.5	4.61943170363406	4.583551188456594

*In test three the result has nothing to do with choosing, but only whether the arm gives a reward since both arms have a probability of 0.5. So it shows that the result of both policies still greatly depends on probability, neither is guaranteed to yield the best result.

*6. reference reading:

- 1 Richard S. Sutton and Andrew G. Barto, *Reinforcement Learning: An Introduction*
- 2 Lucas Descauses, *Reinforcement Learning with Function Approximation in Continuing Tasks: Discounted Return or Average Reward?*